

A Comparative Study of Locking Techniques: Fairness and Throughput Performance Evaluation

Ching-Shen Lin, Chao-Ting Lin and Shiwu Lo

1

Department of Computer Science
National Chung Cheng University

Outline

- **Introduction**
- **Synchronization Algorithms**
- **Testing Program**
- **Throughput Evaluation**
- **Fairness Evaluation**
- **Conclusion**

3

Introduction

Background

- Synchronization
- GNU/Linux's two common locking strategies
 - Spin Locks
 - Blocking Locks
- **Goal:** Examines the **throughput** and **fairness** of different locking methods in a shared resource access scenario

5

Synchronization Algorithms

Synchronization Algorithms

- Raw spinlock
- Pthread spinlock
- Ticket spinlock
- MCS spinlock
- Pthread Mutex
- Binary Semaphore

Raw Spinlock

- The Raw Spinlock algorithm employs the **test-and-test-and-set (TTAS)** technique to continuously examine and establish the lock status.
- This mechanism guarantees that **only one thread** can have exclusive access to the critical section at any given moment.

Pthread Spinlock

- The GNU Pthreads spinlocks (plock) provided by C library
- It also uses TTAS

```
2 pthread_spinlock_t spin;
3
4 void spin_init(){
5     pthread_spin_init(&spin, PTHREAD_PROCESS_PRIVATE);
6 }
7 void spin_lock(){
8     pthread_spin_lock(&spin);
9 }
10 void spin_unlock(){
11     pthread_spin_unlock(&spin);
12 }
```

Ticket Spinlock

- Threads wait until the "service number" equals their ticket number before entering the critical section.
- This method can consume significant Network- on-Chip (NoC) bandwidth as waiting threads continuously query the "grant" variable in a loop.

MCS Spinlock

- The MCS spinlock algorithm queues waiting threads in a **linked list**.
- Upon leaving the critical section, a thread sets the “locked” variable of the **next thread to false** to signal that it can enter the critical section.

Pthread Mutex

- We used the GNU/Linux's mutex function from the GNU pthread library.
- If the mutex is already locked, the calling thread will be blocked until it acquires the mutex.
- Upon completion of the function, the mutex object will be locked and owned by the calling thread

Binary Semaphore

- We used a binary semaphore with C library to restrict access to one thread at a time.
- `sem_wait()`: decreases the semaphore value if greater than zero, blocking the thread otherwise.
- `sem_post()`: reduces the semaphore by one and resumes waiting threads when the value increases.

13

Testing Program

Testing Environment

- Model name: AMD Ryzen Threadripper 2990WX
- Number of Cores: 32-Core Processor
- Virtual core: 64 (virtual core per core is 2)
- Architecture: x86_64
- Compile Environment: gcc version 10.3.0
- Target: AMD Zen+ and x86 64-linux-gnu
- Thread model: POSIX

Testing Program

- We analyze each lock method in a quantitative manner through a controlled microbenchmark.
- Each thread is assigned to a specific core for controlled testing.

Testing Program

Algorithm 1 The Microbenchmark

```
1  while(1):
2      spin_lock();
3      // critical section start
4      for(each element in sharedData):
5          if(with file I/O):
6              element += 1;
7              fprintf(test_file);
8          if(without file I/O):
9              element += 1;
10     // end of critical section
11     spin_unlock();
12     // remainder section
13     rs_start = clock_gettime();
14     do{
15         t = clock_gettime();
16     }while(t-rs_start < RS_size*rand(0.85~1.15));
```

- The benchmark is run in two variations: one **with file input/output** operations and one **without**.

Testing Program

Algorithm 1 The Microbenchmark

```
1  while(1):
2      spin_lock();
3      // critical section start
4      for(each element in sharedData):
5          if(with file I/O):
6              element += 1;
7              fprintf(test_file);
8          if(without file I/O):
9              element += 1;
10     // end of critcal section
11     spin unlock();
12     // remainder section
13     rs_start = clock_gettime();
14     do{
15         t = clock_gettime();
16     }while(t-rs_start < RS_size*rand(0.85~1.15));
```

- The random variable is used to simulate different loads of the program.

18

Throughput Evaluation

Throughput Evaluation Aspects

- Critical section size
- Remainder section size
- Oversubscription issue

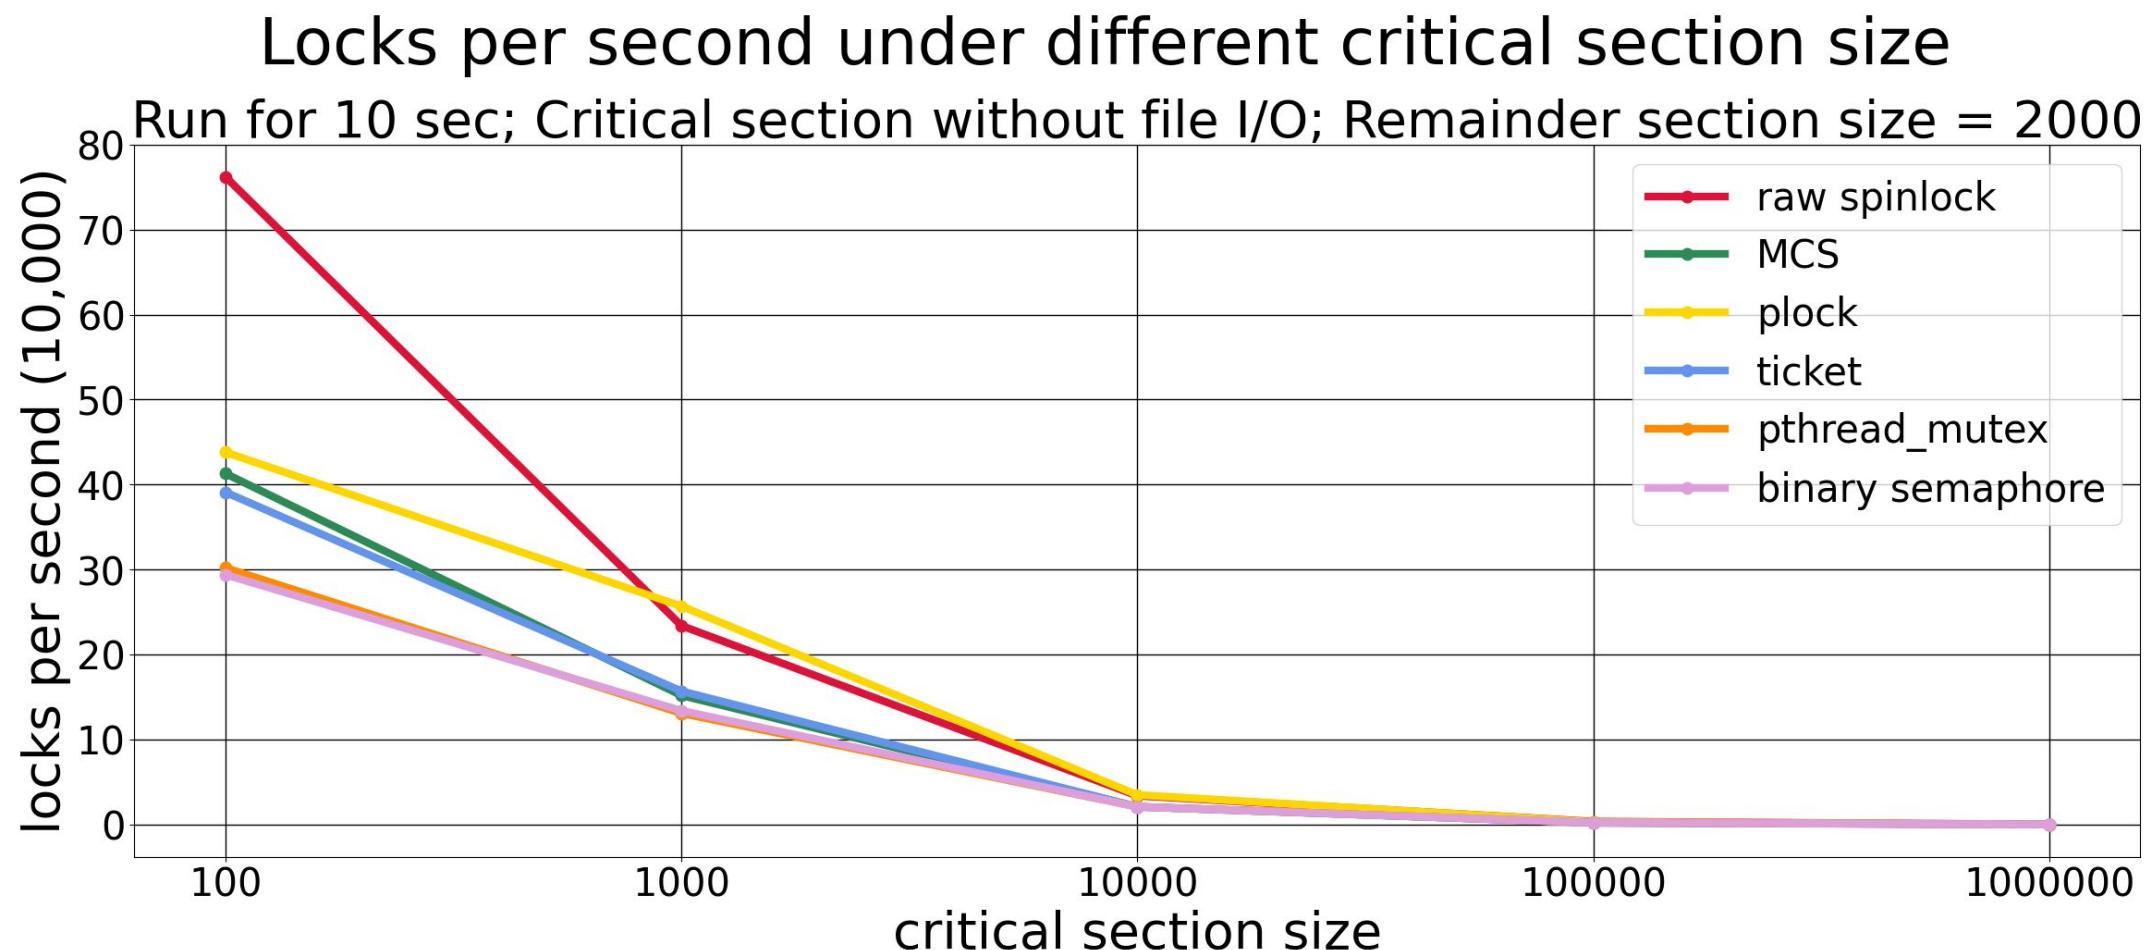
Evaluation Aspects – (1)

- Critical section size

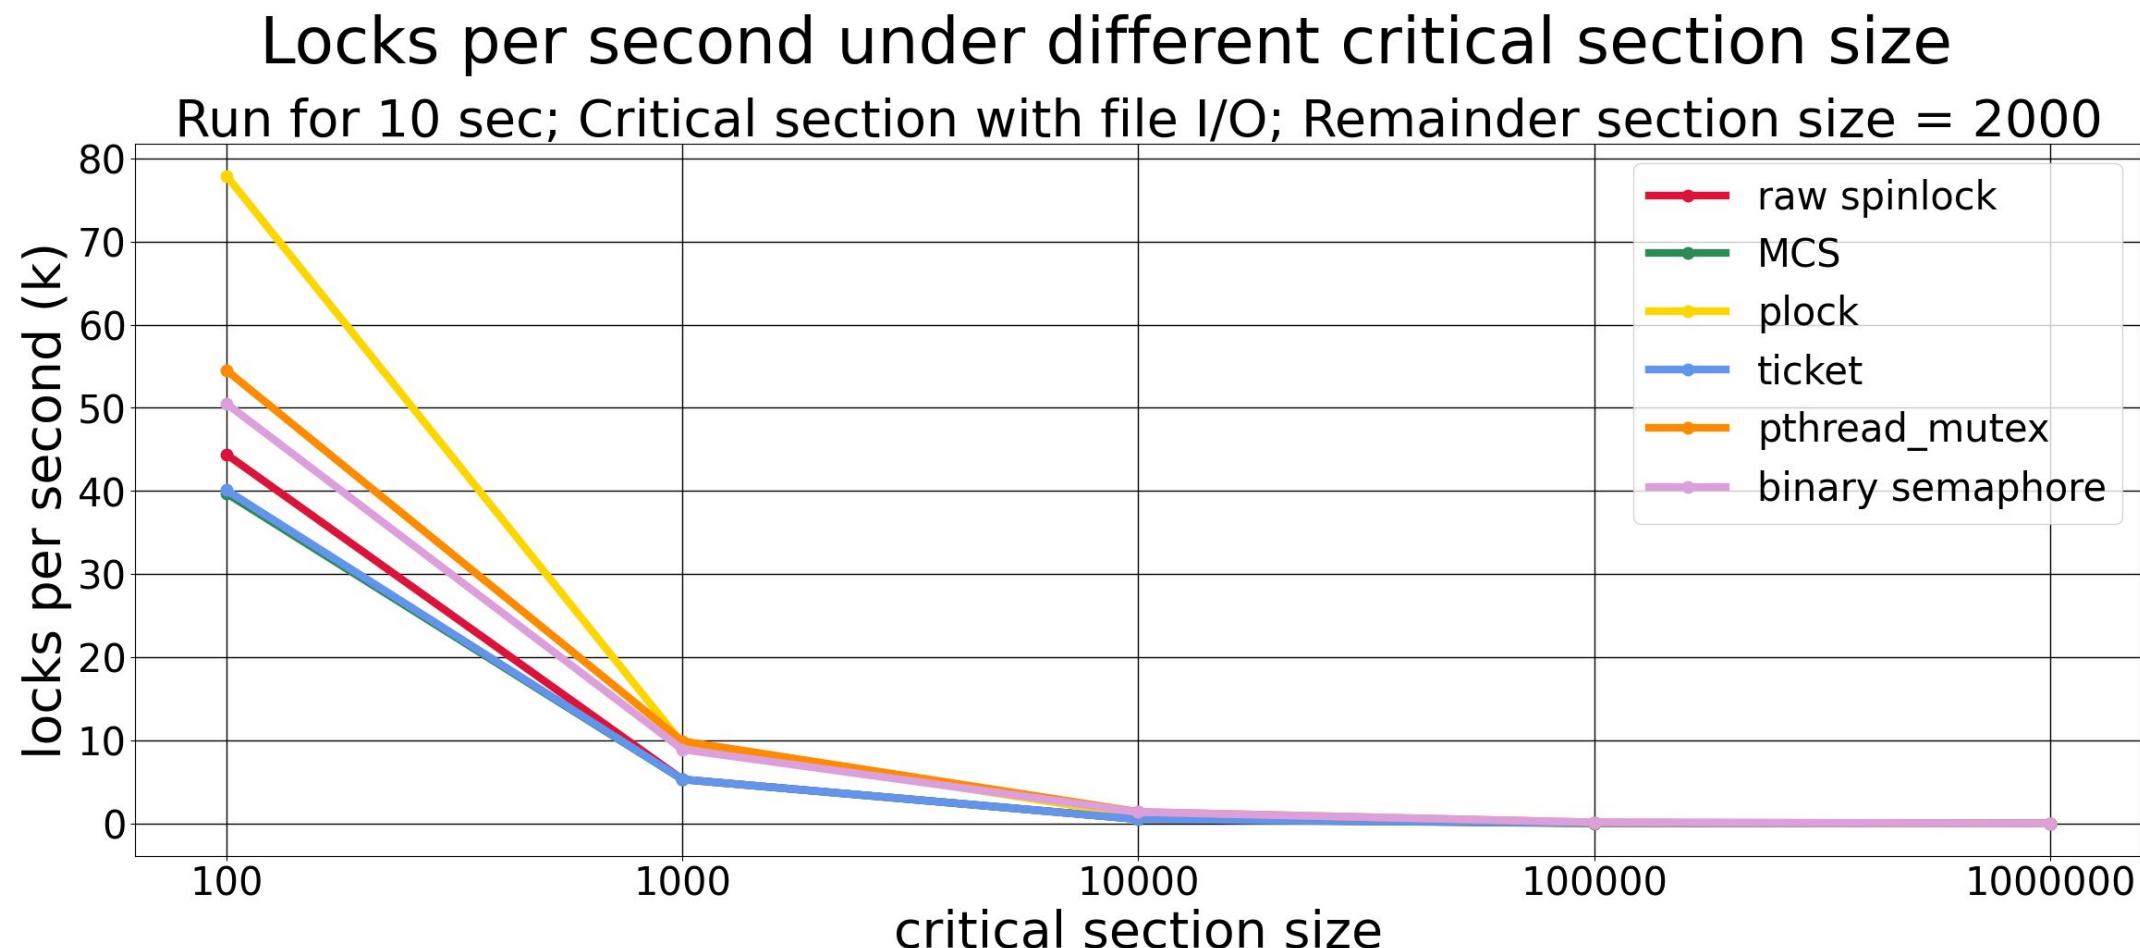
- The number of times the *for* loop execute in the critical section
- For example *for(0 to 100)* means the critical section size = 100.

```
3      // critical section start
4      for(each element in sharedData):
5          if(with file I/O):
6              element += 1;
7              fprintf(test_file);
8          if(without file I/O):
9              element += 1;
10         // end of critcal section
```

Result - Critical Section Size (1) without file I/O



Result - Critical Section Size (2) with file I/O



Evaluation Aspects – (2)

- **Remainder section size**

- The time threads will **wait after leaving the critical section**
- A **larger remainder section size** results in longer waiting times, indicating a **low load condition**. Conversely, a smaller remainder section size corresponds to a high load condition.

```
12      // remainder section
13      rs_start = clock_gettime();
14      do{
15          t = clock_gettime();
16      }while(t-rs_start < RS_size*rand(0.85~1.15));
```

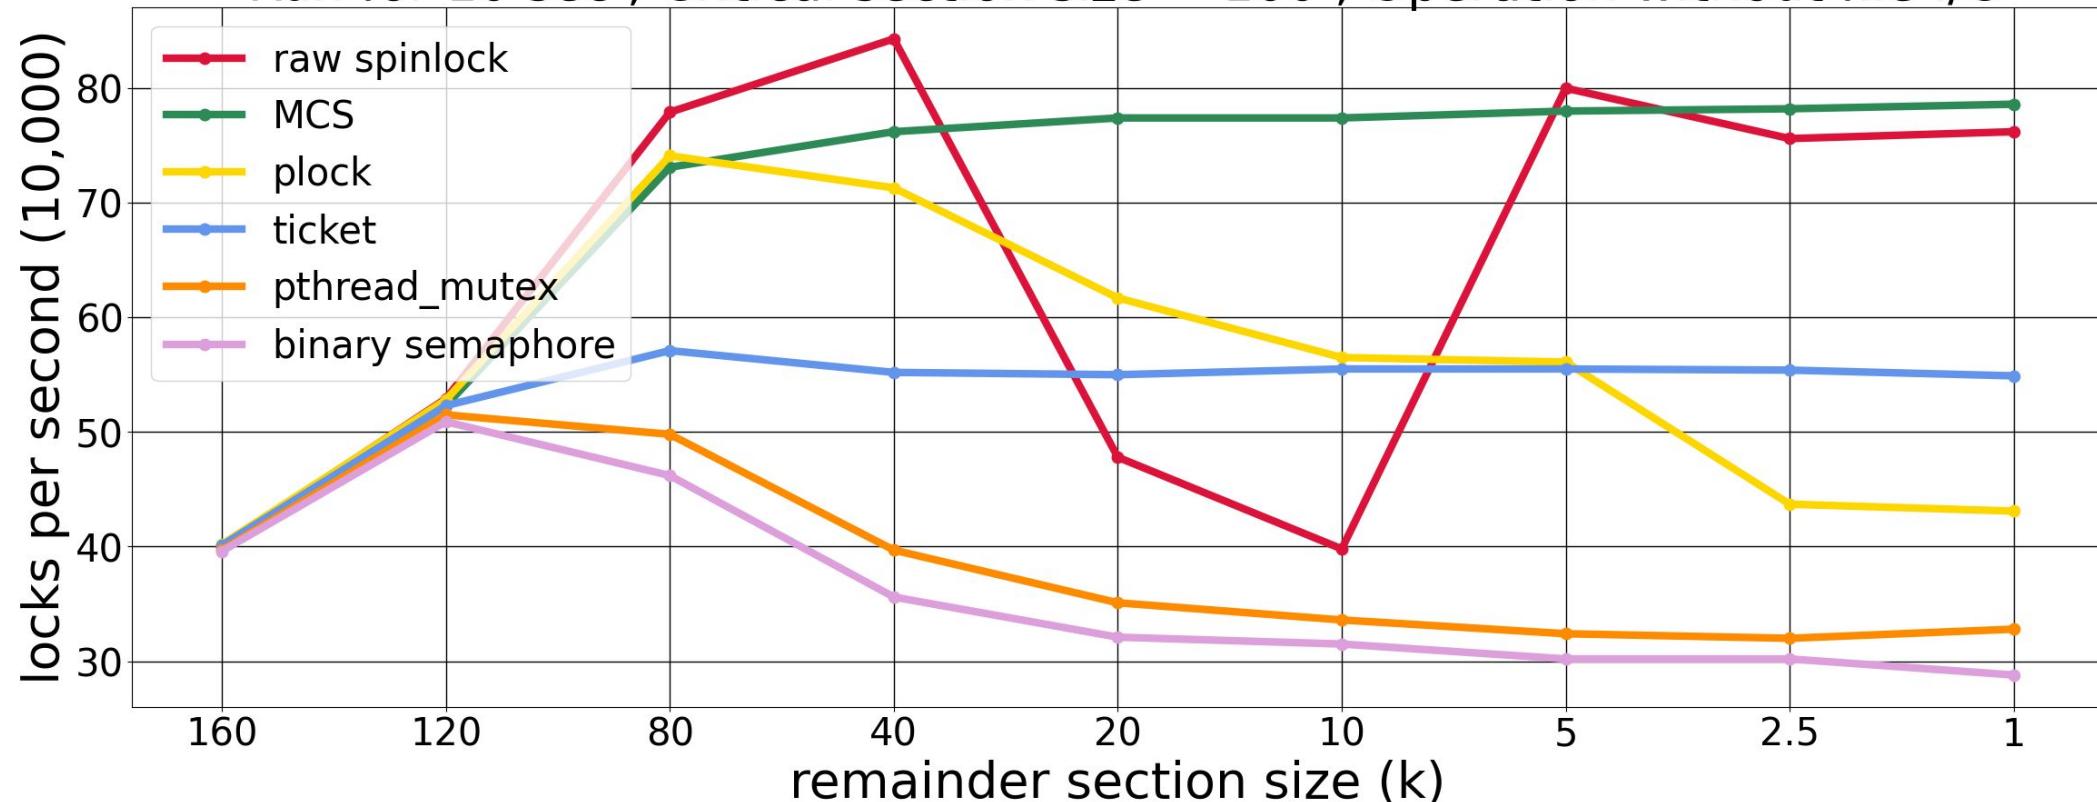
Result - Remainder Section Size

(1)

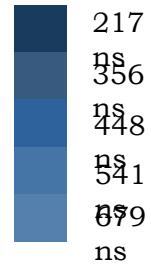
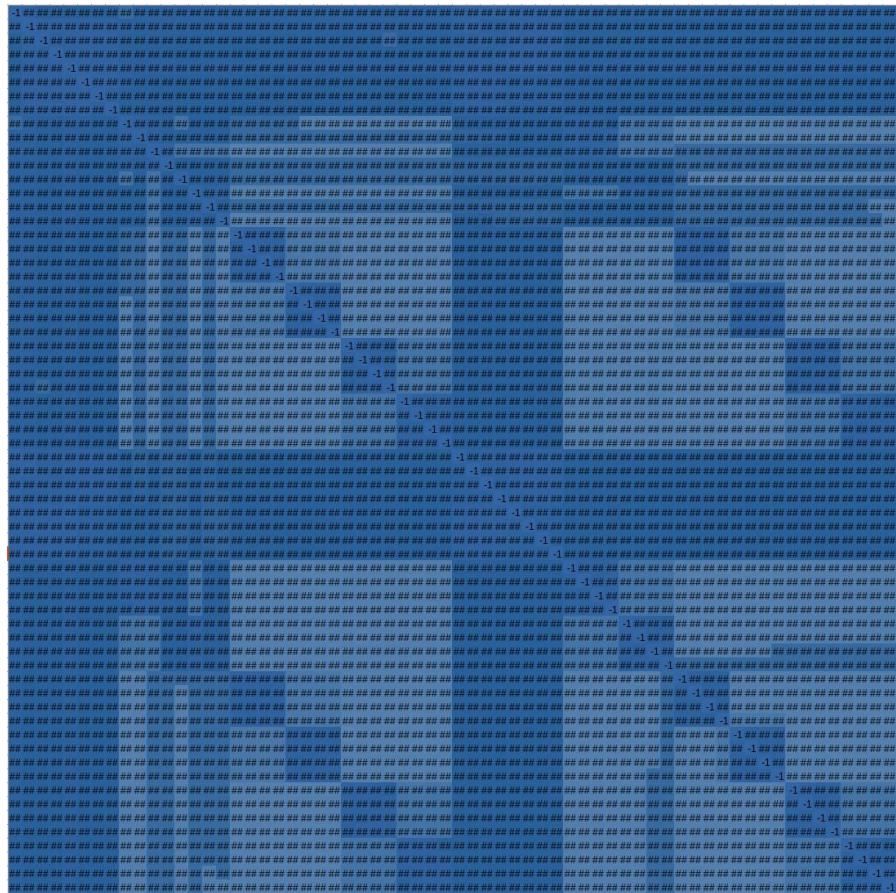
without file I/O

Locks per second under different remainder section size

Run for 10 sec ; Critical section size = 100 ; Operation without file I/O



Why Raw Spinlock drops



- This graph shows the transmission time between cores
- The deeper the closer between two cores
- We use *atomic_compare_exchange()*
- Under this workload, the cores involved in the transfer happen to be farther apart, resulting in decreased performance

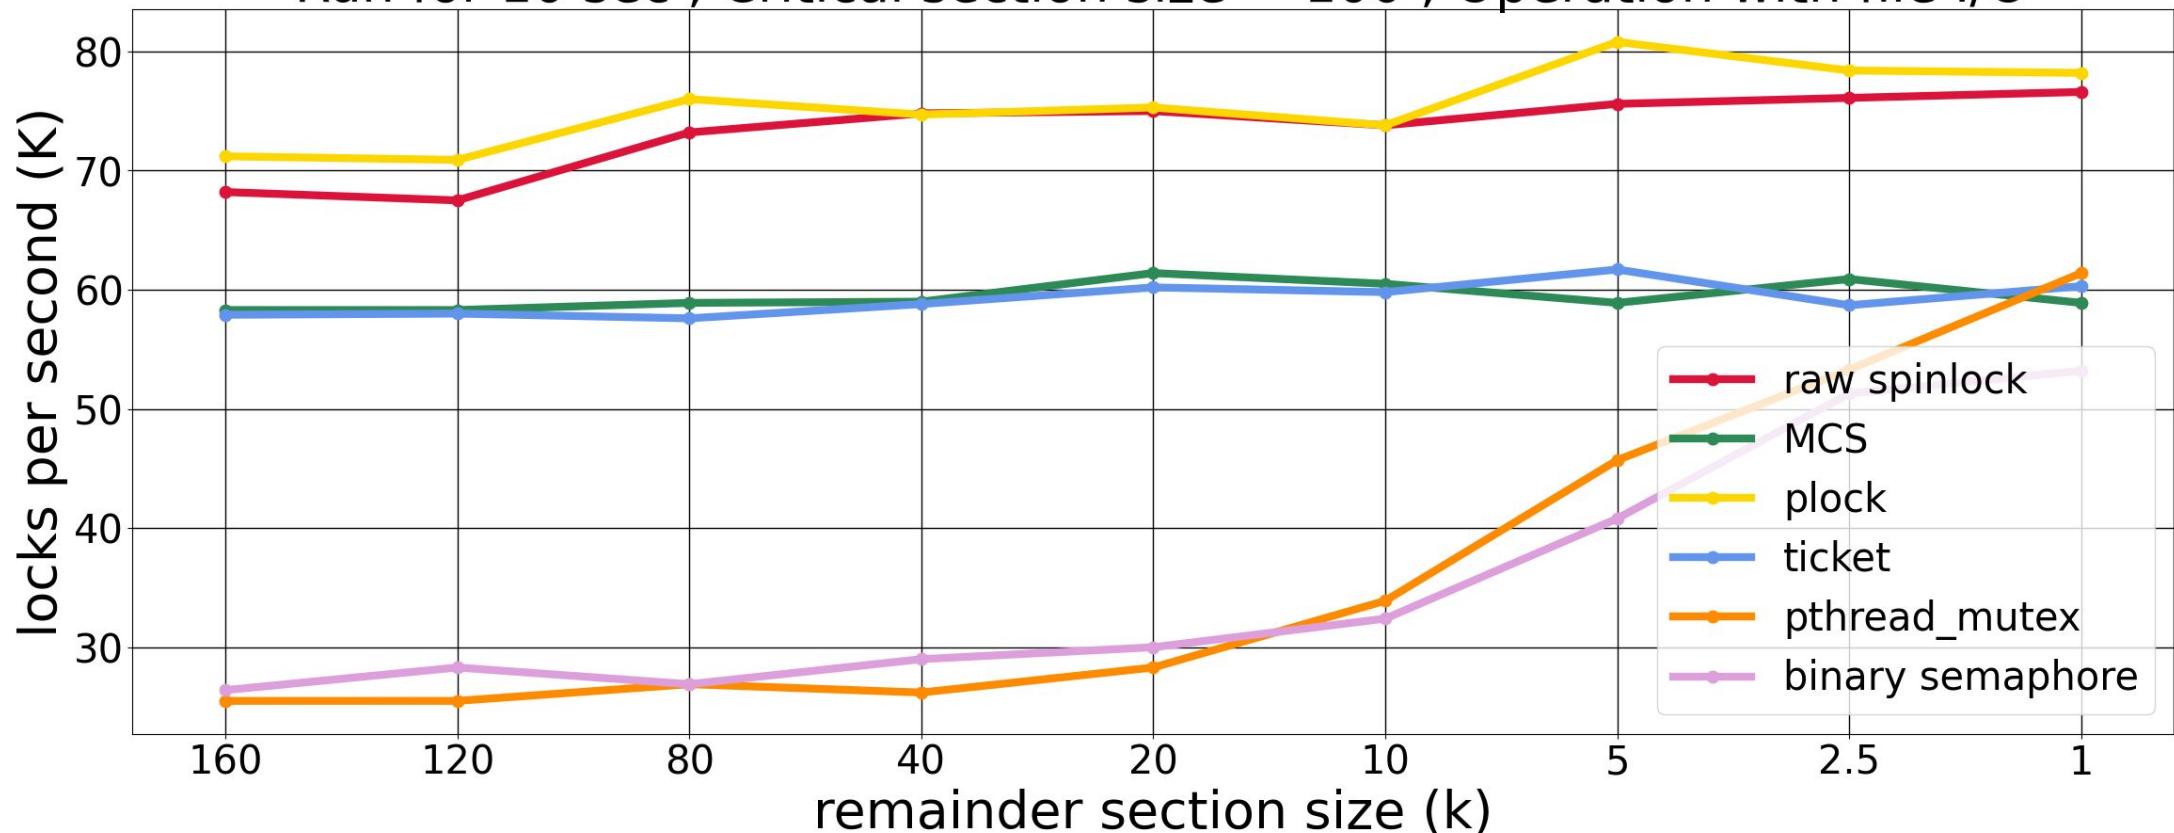
Result - Remainder Section Size

(2)

with file I/O

Locks per second under different remainder section size

Run for 10 sec ; Critical section size = 100 ; Operation with file I/O



Evaluation Aspects – (3)

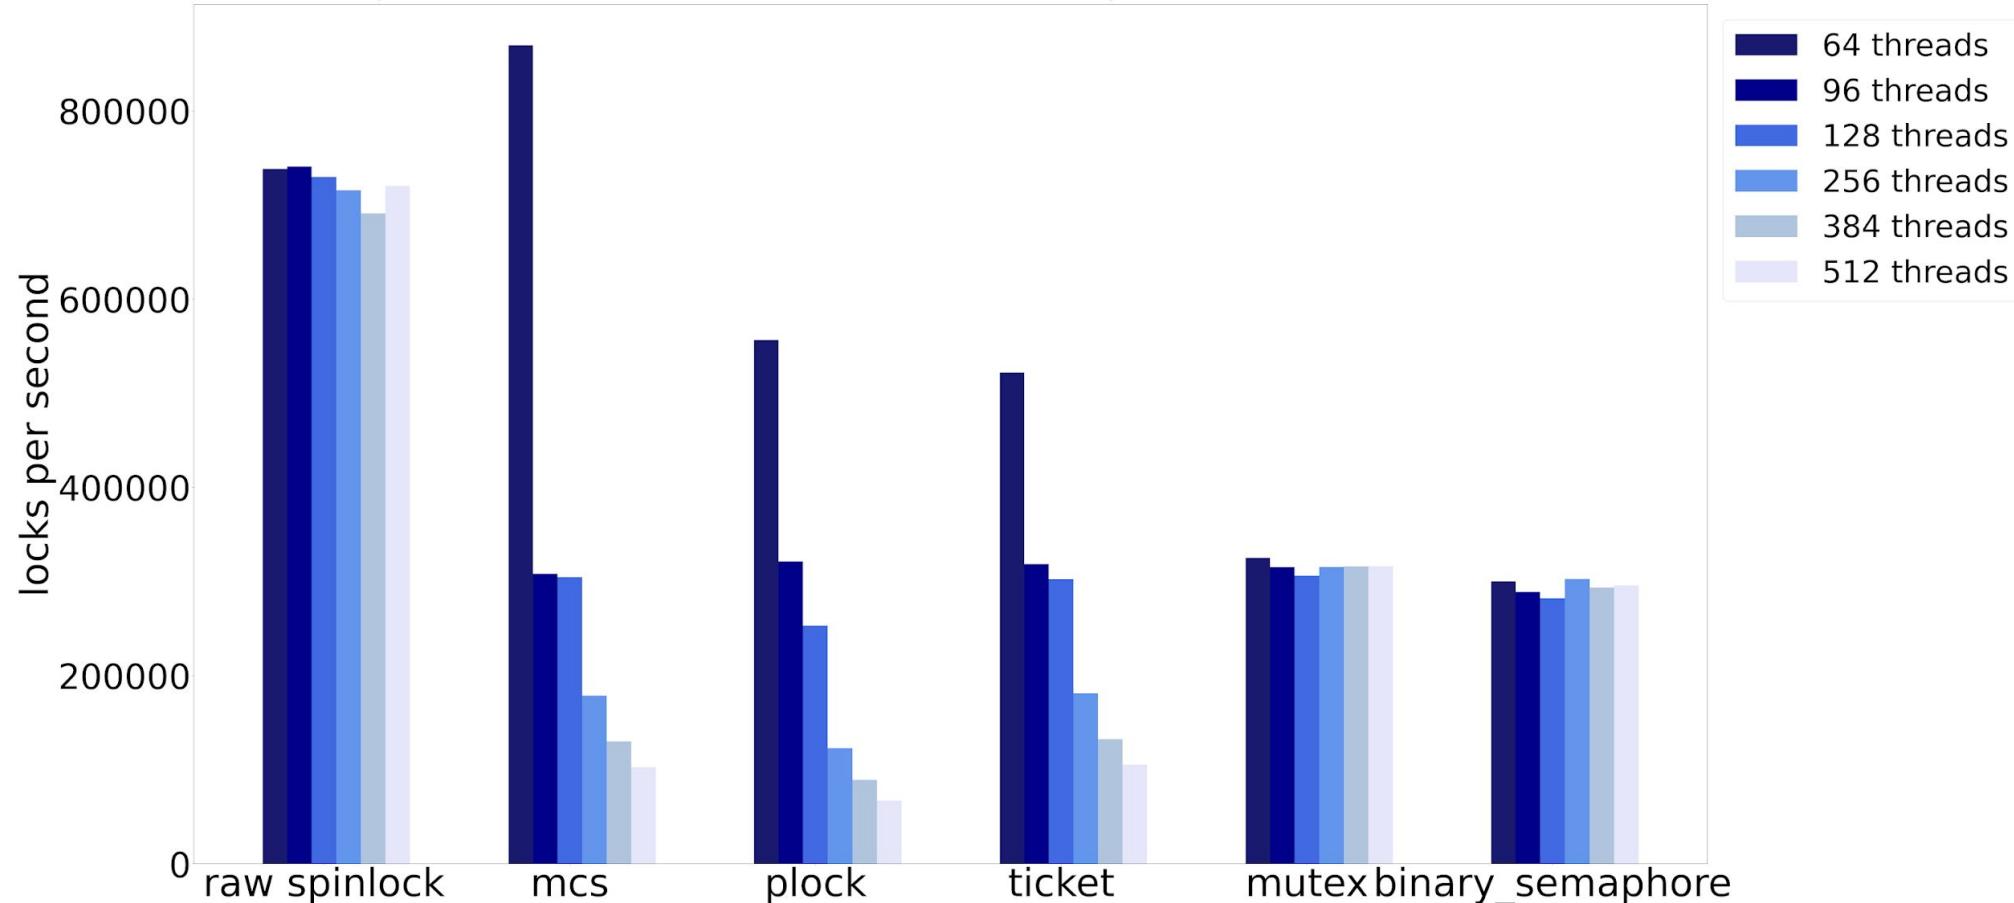
- **Oversubscription Issue**

- In oversubscription, when there are **more threads than available cores**, it causes waiting, **increased latency, and decreased performance**.
- We evaluated the locks for oversubscription by running the microbenchmark on a system **with an excess of threads compared to cores**, measuring their performance.

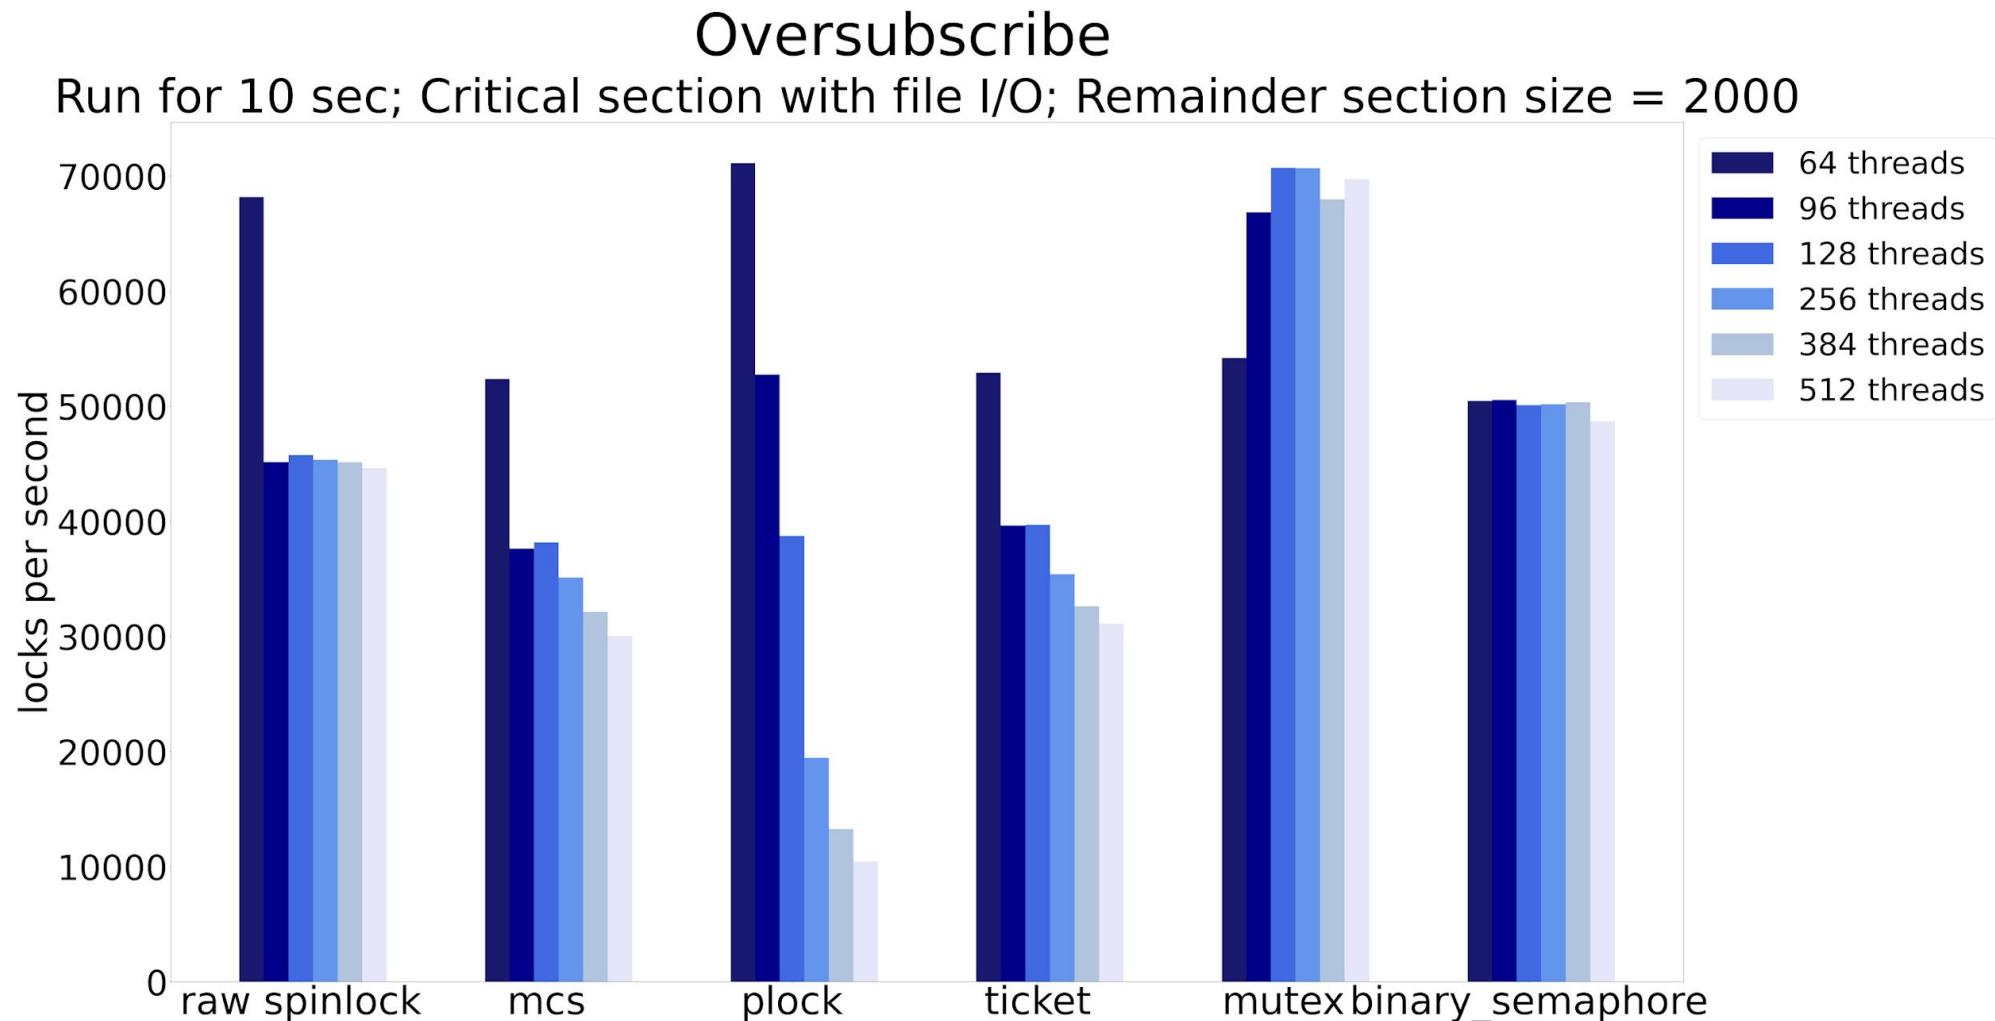
Result - Oversubscription Issue (1) without file I/O

Oversubscribe

Run for 10 sec; Critical section without file I/O; Remainder section size = 2000



Result - Oversubscription Issue (2) with file I/O



30

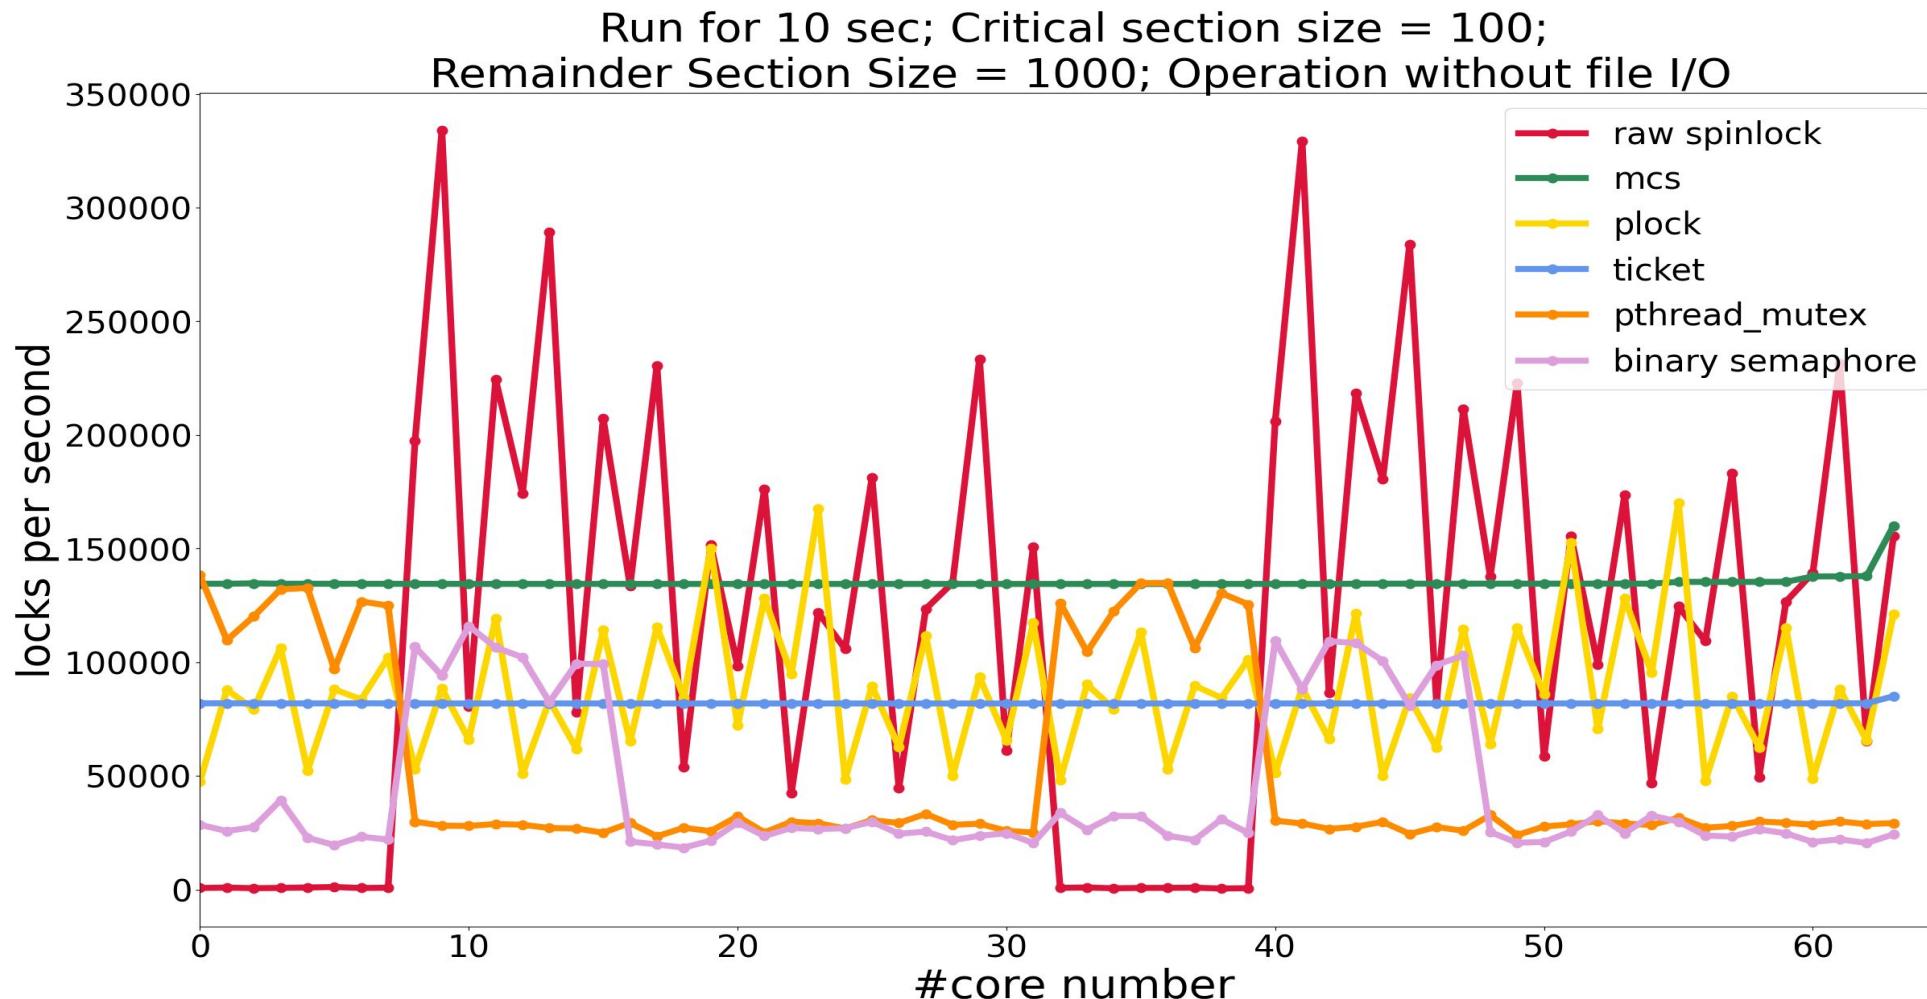
Fairness Evaluation

Result - Undersubscription Issue

(1)

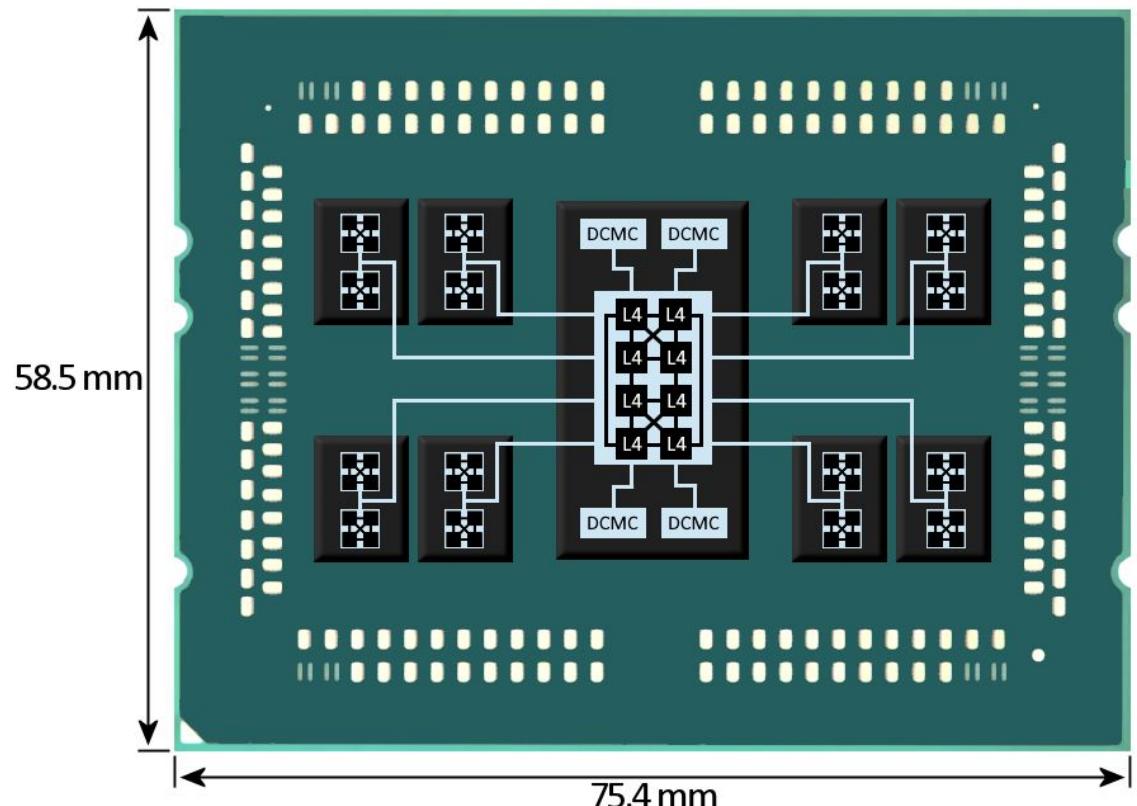
without file I/O

Fairness Evaluation

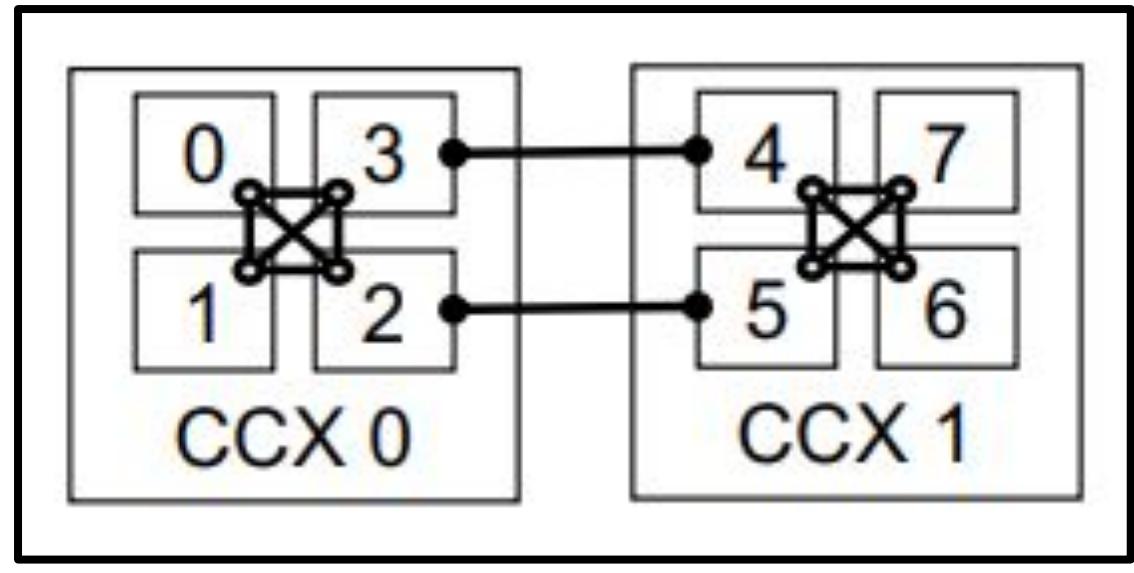


Why

AMD EPYC Rome (topology speculation)



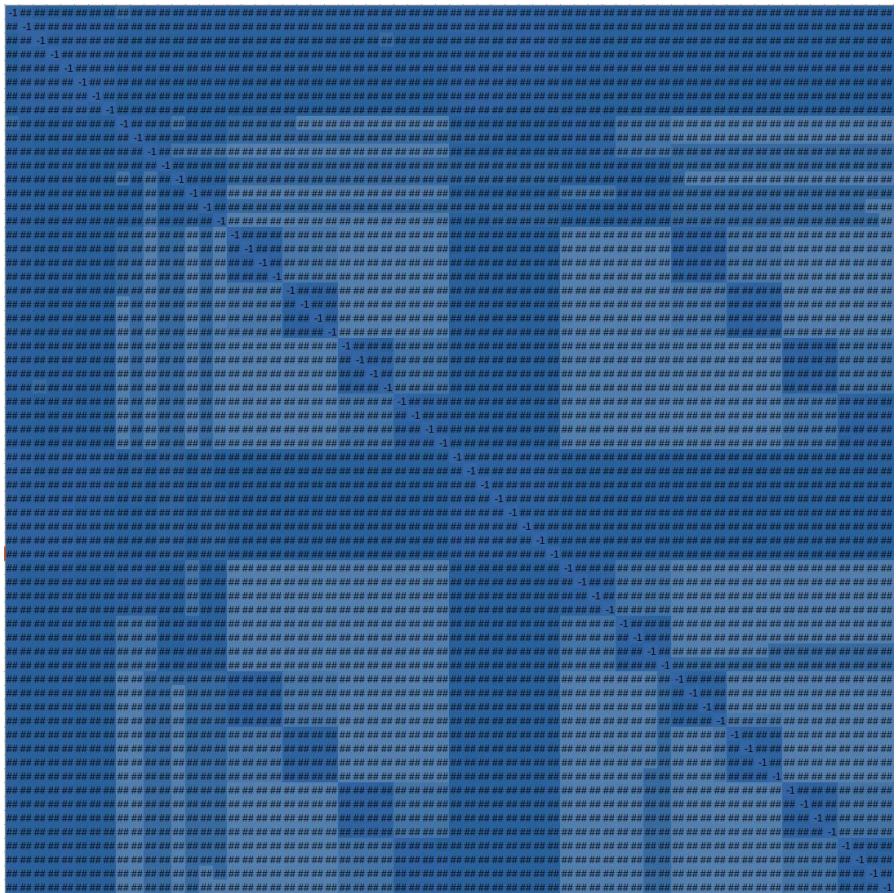
AMD EPYC processor architecture



CPU Die

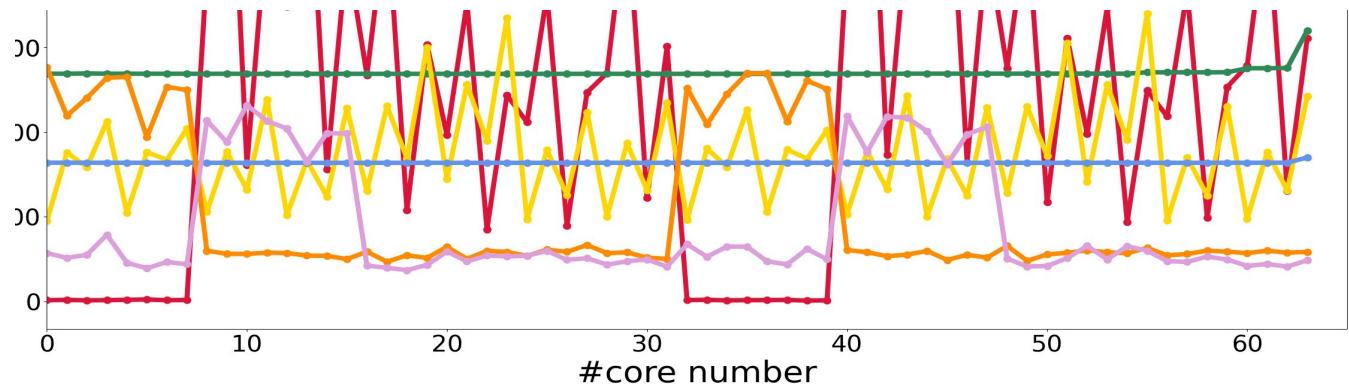
Resource

Why



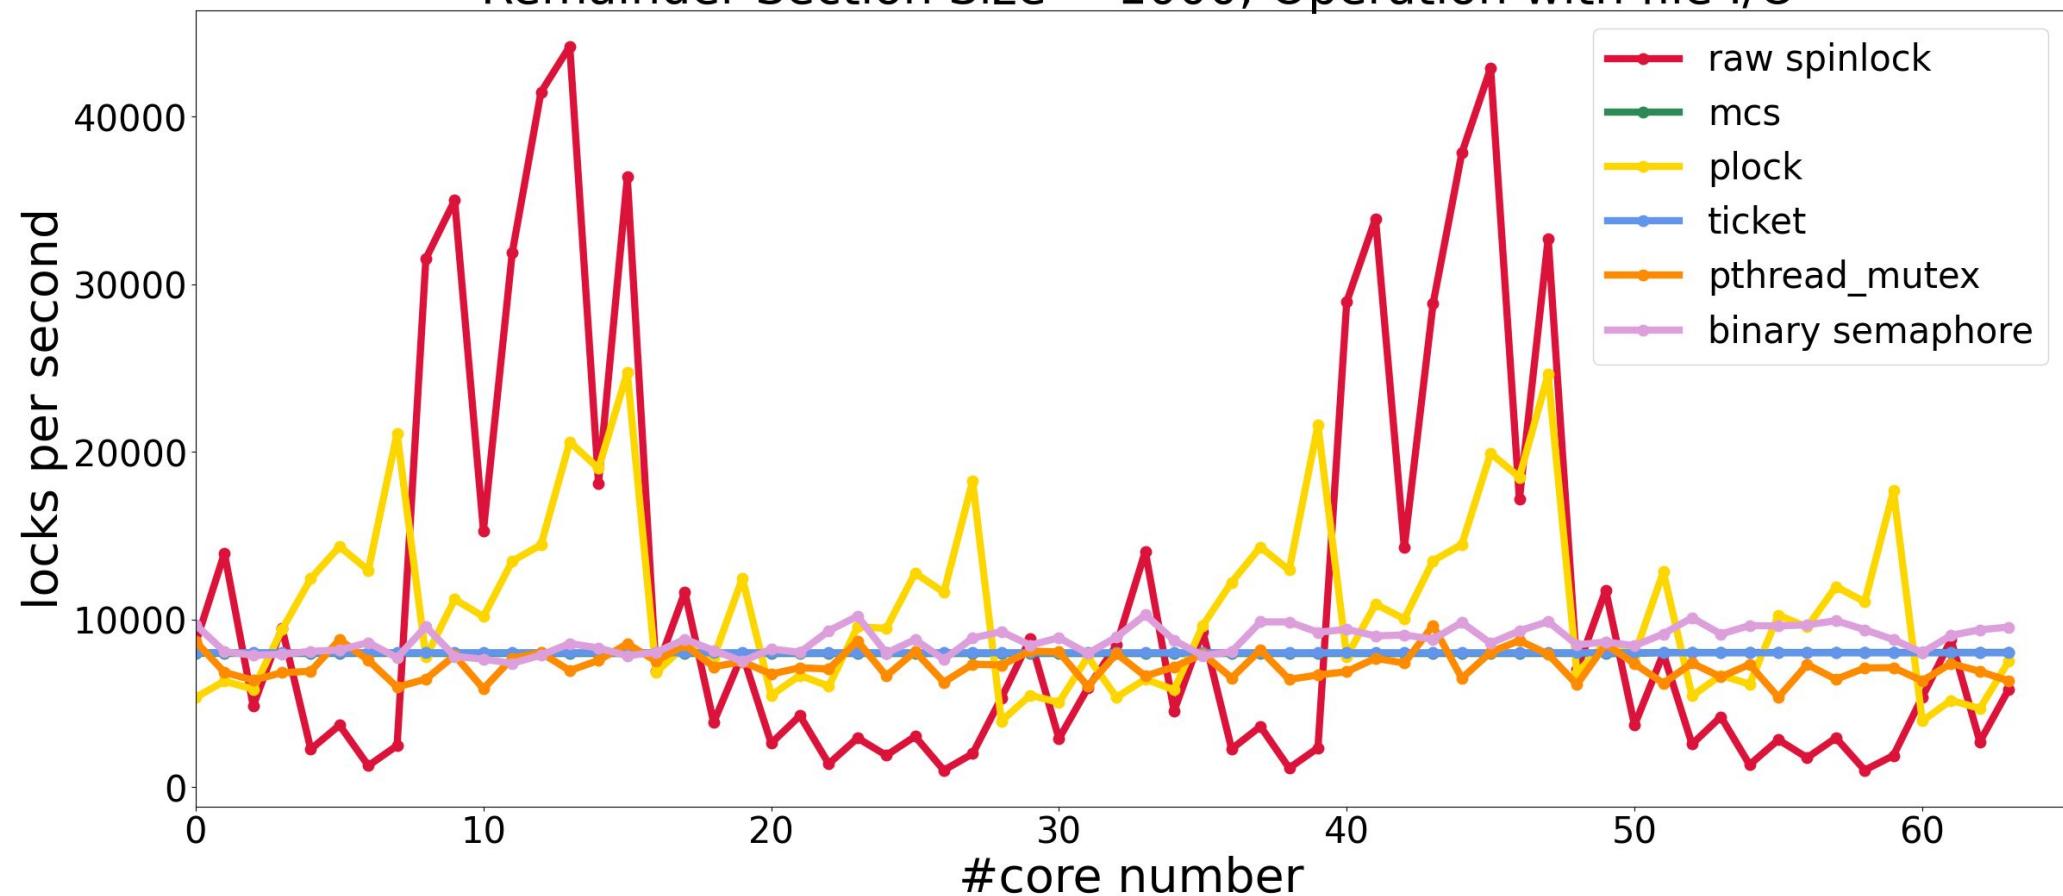
217
ns
356
ns
448
ns
541
ns
679
ns

- This graph shows the transmission time between cores
- The deeper the closer between two cores
- The distances between cores can be divided into several major categories.

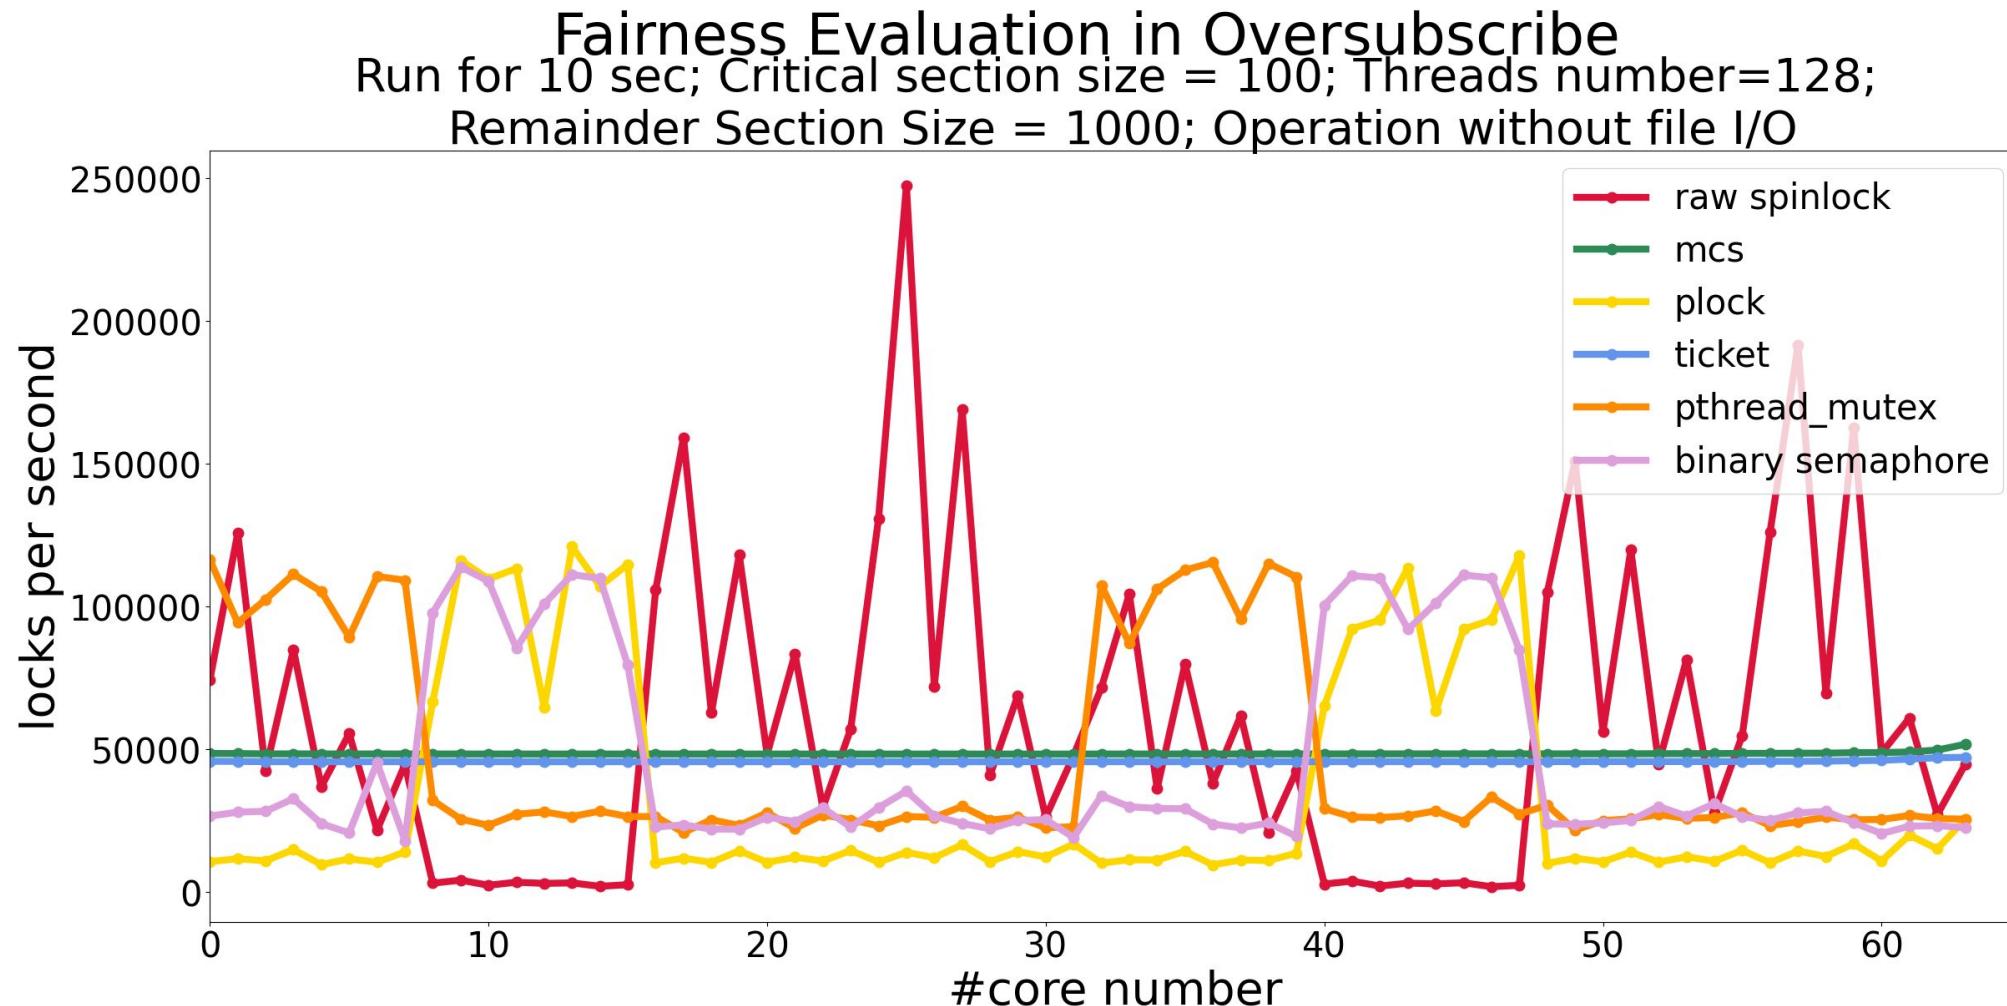


Result - Undersubscription Issue (2) with file I/O

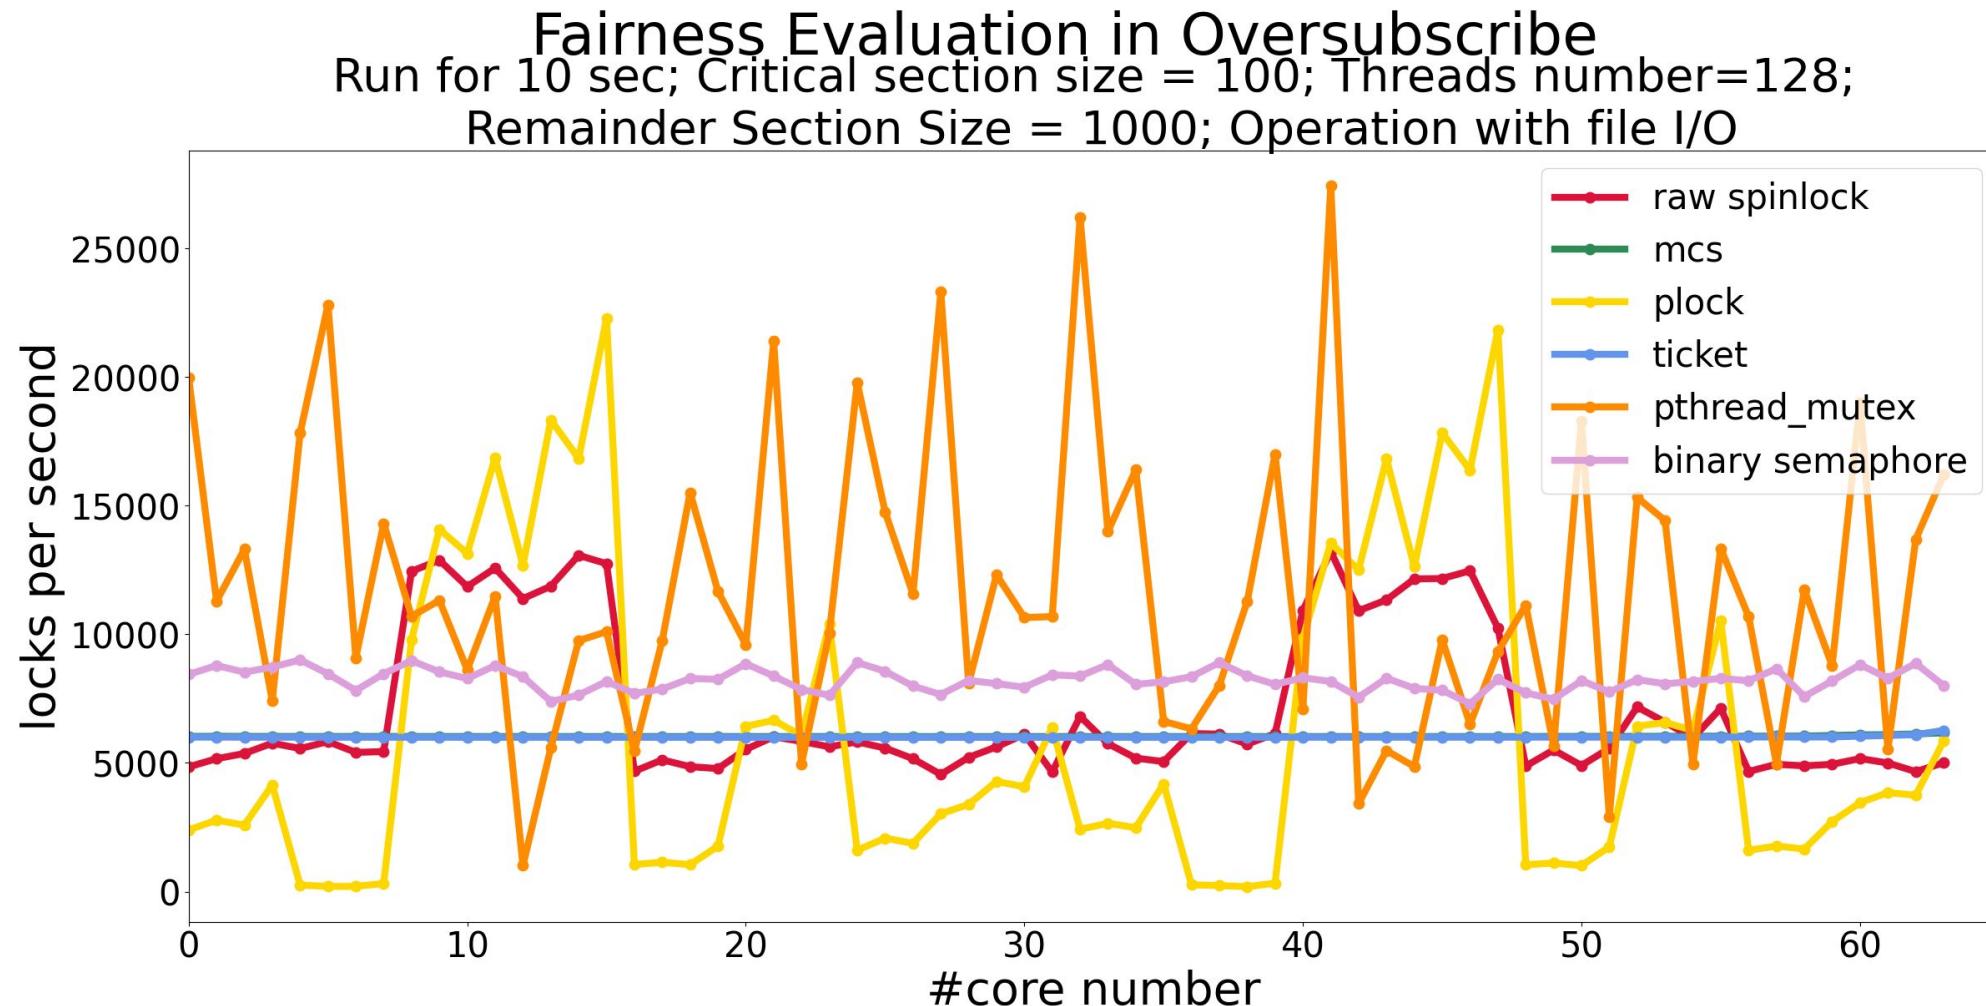
Fairness Evaluation
Run for 10 sec; Critical section size = 100;
Remainder Section Size = 1000; Operation with file I/O



Result - Oversubscription Issue (1) without file I/O



Result - Oversubscription Issue (2)



Conclusion

- Locking mechanism choice is critical for performance and fairness in shared resource access.
- Spinlocks perform well in smaller critical section sizes.
- Semaphore and mutex locks outperform ticket and MCS spinlocks as critical section size increases.
- The type of operation performed on the shared resource affects locking mechanism performance. Mutex and semaphore locks excel in file I/O operations.

Conclusion

- Oversubscription negatively impacts performance and fairness across all locking mechanisms.
- GNU/Linux's mutex exhibits extreme unfairness during file I/O operations in oversubscription.
- Optimal performance and fairness in shared resource access depend on carefully considering critical section size, operation type, and the potential for oversubscription when selecting a locking mechanism.

39

Q&A

40

**Thank you for your
attention**